

GitHub: <https://github.com/sgudishetty9/cloudcomputing-proj2>

DockerHub:

Repo: <https://hub.docker.com/repository/docker/sg2489/wqprepo/general>

Image:

<https://hub.docker.com/layers/sg2489/wqprepo/v1/images/sha256-9416871c59604d5027cbbaf0e2060cf0d6661dc1a891e882939b60d63f46f08c?context=repo>

Procedure:

Cluster Creation


Navigate to Amazon EMR in AWS > Create the Cluster, and select the following apps:


Name and applications [Info](#)


Name


Amazon EMR release [Info](#)
A release contains a set of applications which can be installed on your cluster.


Application bundle


Spark
Interactive



Core
Hadoop


Flink


HBase


Presto


Trino


Custom


☐ Flink 1.17.1
☐ HCatalog 3.1.3
☐ Hue 4.11.0
☐ Livy 0.7.1
☐ Phoenix 5.1.3
☒ Spark 3.4.1
☐ Tez 0.10.2
☐ ZooKeeper 3.5.10

☐ Ganglia 3.7.2
☒ Hadoop 3.3.6
☐ JupyterEnterpriseGateway 2.6.0
☐ MXNet 1.9.1
☐ Pig 0.17.0
☐ Sqoop 1.4.7
☐ Trino 426

☐ HBase 2.4.17
☐ Hive 3.1.3
☐ JupyterHub 1.5.0
☐ Oozie 5.2.1
☐ Presto 0.283
☐ TensorFlow 2.11.0
☒ Zeppelin 0.10.1

Had to train model using 4 EC2 instances, so configured 'Provisioning Configuration as below:

Cluster scaling and provisioning
[Info](#)

Set up scaling and provisioning configurations for the core and task node groups for your cluster.

Choose an option

☒ **Set cluster size manually**
 Use this option if you know your workload patterns in advance.

☐ **Use EMR-managed scaling**
 Monitor key workload metrics so that EMR can optimize the cluster size and resource utilization.

☐ **Use custom automatic scaling**
 To programmatically scale core and task nodes, create custom automatic scaling policies.

Provisioning configuration

Set the size of your core and task instance groups. Amazon EMR attempts to provision this capacity when you launch your cluster.

Name	Instance type	Instance(s) size	Use Spot purchasing option
Core	m5.xlarge	<input type="text" value="1"/>	<input type="checkbox"/>
Task - 1	m5.xlarge	<input type="text" value="3"/>	<input type="checkbox"/>

Went with Recommended Cluster Termination settings:

Cluster termination
[Info](#)

☐ Manually terminate cluster
☐ Automatically terminate cluster after last step ends
☒ **Automatically terminate cluster after idle time (Recommended)**

Idle time

Enter the time until your cluster terminates.

Choose a time that is greater than 1 minute (00:01:00) and less than 7 days. The time is in hh:mm:ss (24-hour) format.

☒ **Use termination protection**
 Protect your EC2 instances from accidental termination.

Create a key pair for SSH and SCP to the Instance/Cluster:

Security configuration and EC2 key pair - optional
[Info](#)

Security configuration

Select your cluster encryption, authentication, authorization, and instance metadata service settings.

Amazon EC2 key pair for SSH to the cluster [Info](#)

Select default IAM roles:

Identity and Access Management (IAM) roles [Info](#)
Choose or create a service role and instance profile for the EC2 instances in your cluster.

Amazon EMR service role [Info](#)
The service role is an IAM role that Amazon EMR assumes to provision resources and perform service-level actions with other AWS services.

☒ **Choose an existing service role**
Select a default service role or a custom role with IAM policies attached so that your cluster can interact with other AWS services.

☐ **Create a service role**
Let Amazon EMR create a new service role so that you can grant and restrict access to resources in other AWS services.

Service role

EMR_DefaultRole

EC2 instance profile for Amazon EMR
The instance profile assigns a role to every EC2 instance in a cluster. The instance profile must specify a role that can access the resources for your steps and bootstrap actions.

☒ **Choose an existing instance profile**
Select a default role or a custom instance profile with IAM policies attached so that your cluster can interact with your resources in Amazon S3.

☐ **Create an instance profile**
Let Amazon EMR create a new instance profile so that you can specify a custom set of resources for it to access in Amazon S3.

Instance profile

EMR_EC2_DefaultRole

Rest, go with default configuration, cluster/s will be working upon turning to 'Waiting' status.

WQP Updated less than a minute ago [Refresh](#) [Terminate](#) [Clone in AWS CLI](#) [Clone](#)

▼ Summary

Cluster info	Applications	Cluster management	Status and time
Cluster ID j-2Y5VVES1XCOTB	Amazon EMR version emr-6.15.0	Log destination in Amazon S3 aws-logs-098431872762-us-east-1/elasticmapreduce	Status ✔ Waiting
Cluster configuration Instance groups	Installed applications Hadoop 3.3.6, Spark 3.4.1, Zeppelin 0.10.1	Persistent application UIs Spark History Server YARN timeline server	Creation time December 08, 2023, 10:29 (UTC-05:00)
Capacity 1 Primary 1 Core 3 Task		Primary node public DNS ec2-3-234-217-212.compute-1.amazonaws.com Connect to the Primary node using SSH Connect to the Primary node using SSM	Elapsed time 13 minutes, 41 seconds

Create an S3 Bucket and drop training and testing files into it.

sg2489-wqp
Info

Objects
Properties
Permissions
Metrics
Management
Access Points

Objects (3)
Info

Objects are the fundamental entities stored in Amazon S3. You can use [Amazon S3 Inventory](#) to get a list of all objects in your bucket. For others to access your objects, you'll need to explicitly grant them permissions. [Learn more](#)

Refresh
Copy S3 URI
Copy URL
Download
Open
Delete
Actions
Create folder
Upload

Find objects by prefix

	Name	Type	Last modified	Size	Storage class
<input type="checkbox"/>	TrainingDataset.csv	csv	December 8, 2023, 10:39:54 (UTC-05:00)	67.2 KB	Standard
<input type="checkbox"/>	ValidationDataset.csv	csv	December 8, 2023, 10:39:55 (UTC-05:00)	8.6 KB	Standard
<input type="checkbox"/>	wine_quality_prediction_train.py	py	December 8, 2023, 10:39:54 (UTC-05:00)	3.9 KB	Standard

Ensure master instance of the cluster has permission to SSH, if not edit inbound rule after navigating to the security group of the instance.

Inbound rules (8)
Manage tags
Edit inbound rules

Search

1

	Security group rule...	IP version	Type	Protocol	Port range	Source
7	sgr-0aa323dbe1b6c30...	IPv4	SSH	TCP	22	96.24...

Done SSH into my master instance using below command:
ssh -i "<my-key-pair.pem>" ec2-user@<Public DNS>

Fetch trained model from S3 to master instance into a new folder (trained_model) using below command:

```
aws s3 sync aws s3 sync s3://<my-bucket-name>/trained_model.model ./trained_model
```

Done SCP for <testdataset> and Single Machine Prediction code file using below command:

```
scp -i <keypair.pem> <file> ec2-user@<Public DNS>:/home/ec2-user
```

Installed Pyspark with: pip install pyspark

for running single machine prediction application with external test dataset, using below command:

```
python wqp_single_machine.py
```

```
Sample Predictions:
+-----+
|fixed acidity|volatile acidity|citric acid|residual sugar|chlorides|free sulfur dioxide|total sulfur dioxide|density| pH|sulphates|alcohol|quality|          features|label|    rawPrediction|
+-----+
| 8.9|          0.22|      0.48|      1.8|    0.077|      29.0|      60.0| 0.9968|3.39|    0.53|    9.4|    6.0|[8.9,0.22,0.48,1....]|    1.0|[19.4618621405359...]|0.0389| |
|0372428107...|    1.0|      0.39|      0.31|      2.3|    0.082|      23.0|      71.0| 0.9982|3.52|    0.65|    9.7|    5.0|[7.6,0.39,0.31,2....]|    0.0|[479.531456024878...]|0.9590|
|6291204975...|    0.0|      0.43|      0.21|      1.6|    0.106|      10.0|      37.0| 0.9966|3.17|    0.91|    9.5|    5.0|[7.9,0.43,0.21,1....]|    0.0|[484.885809101067...]|0.9697|
|7161820213...|    0.0|      0.49|      0.11|      2.3|    0.084|      9.0|      67.0| 0.9968|3.17|    0.53|    9.4|    5.0|[8.5,0.49,0.11,2....]|    0.0|[491.163209482913...]|0.9823|
|2641896582...|    0.0|      0.4|      0.14|      2.4|    0.085|      21.0|      40.0| 0.9968|3.43|    0.63|    9.7|    6.0|[6.9,0.4,0.14,2.4...]|    1.0|[3.29191168480642...]|0.0065|
|8382336961...|    1.0|
+-----+
only showing top 5 rows

None
Test Accuracy of the Wine Quality Prediction Model: 0.9843627834245504
/usr/local/lib/python3.7/site-packages/pyspark/sql/context.py:159: FutureWarning: Deprecated in 3.0.0. Use SparkSession.builder.getOrCreate() instead.
FutureWarning,
Weighted F1 score of the Wine Quality Prediction Model: 0.9779339666913879
```

Running Single Machine App in Local with Docker:

Fetch trained model from master instance into local (in a new folder wine_quality_predictor) using below command:

```
scp -i <keypair.pem> -r ec2-user@<Public DNS>:/home/ec2-user/trained_model
./wine_quality_predictor
```

Install and sign up into Docker Desktop, login to Docker from terminal with command: 'docker login'

cd to the source folder and build the image of the app: 'docker build -t <image_name:version> .'

run the app: 'docker run <image_name:version>'

```
Sample Predictions:
+-----+
|fixed acidity|volatile acidity|citric acid|residual sugar|chlorides|free sulfur dioxide|total sulfur dioxide|density| pH|sulphates|alcohol|quality|          features|label|    rawPrediction|
+-----+
| 7.4|          0.7|      0.0|      1.9|    0.076|      11.0|      34.0| 0.9978|3.51|    0.56|    9.4|    5.0|[7.4,0.7,0.0,1.9,...]|    0.0|[488.247586263641...]|0.9764| |
|9517252728...|    0.0|      0.88|      0.0|      2.6|    0.098|      25.0|      67.0| 0.9968|3.2|    0.68|    9.8|    5.0|[7.8,0.88,0.0,2.6...]|    0.0|[476.049269864772...]|0.9520|
|9853972954...|    0.0|      0.76|      0.04|      2.3|    0.092|      15.0|      54.0| 0.997|3.26|    0.65|    9.8|    5.0|[7.8,0.76,0.04,2....]|    0.0|[474.981917756696...]|0.9498|
|0383551339...|    0.0|      0.28|      0.56|      1.9|    0.075|      17.0|      60.0| 0.998|3.16|    0.58|    9.8|    6.0|[11.2,0.28,0.56,1....]|    1.0|[26.3858382129224...]|0.0527|
|7167642584...|    1.0|      0.7|      0.0|      1.9|    0.076|      11.0|      34.0| 0.9978|3.51|    0.56|    9.4|    5.0|[7.4,0.7,0.0,1.9,...]|    0.0|[488.247586263641...]|0.9764|
|9517252728...|    0.0|
+-----+
only showing top 5 rows

None
Test Accuracy of the Wine Quality Prediction Model: 0.96875
Weighted F1 score of the Wine Quality Prediction Model: 0.954791666666667
```

run the app with external dataset: docker run -v
/Users/<username>/<path_to_dataset_directory>:/<path_on_host> <image_name:version>
<external_dataset>

```
-----
Sample Predictions:
-----
|fixed acidity|volatile acidity|citric acid|residual sugar|chlorides|free sulfur dioxide|total sulfur dioxide|density| pH|sulphates|alcohol|quality|      features|label|      rawPrediction|
|probability|prediction|
-----
|      8.9|      0.22|      0.48|      1.8|      0.077|      29.0|      60.0| 0.9968|3.39|      0.53|      9.4|      6.0|[8.9,0.22,0.48,1....|      1.0|[19.4518621405359....|[0.0389
0372428107...|      1.0|
|      7.6|      0.39|      0.31|      2.3|      0.082|      23.0|      71.0| 0.9982|3.52|      0.65|      9.7|      5.0|[7.6,0.39,0.31,2....|      0.0|[479.531456024878....|[0.9590
6291204975...|      0.0|
|      7.9|      0.43|      0.21|      1.6|      0.106|      10.0|      37.0| 0.9966|3.17|      0.91|      9.5|      5.0|[7.9,0.43,0.21,1....|      0.0|[484.885809101067....|[0.9697
7161820213...|      0.0|
|      8.5|      0.49|      0.11|      2.3|      0.084|      9.0|      67.0| 0.9968|3.17|      0.53|      9.4|      5.0|[8.5,0.49,0.11,2....|      0.0|[491.163209482913....|[0.9823
2641896582...|      0.0|
|      6.9|      0.4|      0.14|      2.4|      0.085|      21.0|      40.0| 0.9968|3.43|      0.63|      9.7|      6.0|[6.9,0.4,0.14,2.4....|      1.0|[3.29191168408642....|[0.0065
8382336961...|      1.0|
-----
only showing top 5 rows

None
Test Accuracy of the Wine Quality Prediction Model: 0.9843627834245584
Weighted F1 score of the Wine Quality Prediction Model: 0.9779339666913879
```

Create a new repo in DockerHub.

Push the image into repo with the below commands:

docker tag <img:ver> <user_id>/<repo>:<tag>

docker push <user_id>/<repo>:<tag>

Running Single Machine App in EC2 with Docker:

Create an EC2 instance with default configurations

Install docker with command: sudo yum install docker

Start docker: sudo service docker start

Pull the image into EC2: sudo docker pull <user_id>/<repo>:<tag>

Run the app: sudo docker run <user_id>/<repo>:<tag>

```
-----
Sample Predictions:
-----
|fixed acidity|volatile acidity|citric acid|residual sugar|chlorides|free sulfur dioxide|total sulfur dioxide|density| pH|sulphates|alcohol|quality|      features|label|      rawPrediction|
|probability|prediction|
-----
|      8.9|      0.22|      0.48|      1.8|      0.077|      29.0|      60.0| 0.9968|3.39|      0.53|      9.4|      6.0|[8.9,0.22,0.48,1....|      1.0|[19.4518621405359....|[0.0389
0372428107...|      1.0|
|      7.6|      0.39|      0.31|      2.3|      0.082|      23.0|      71.0| 0.9982|3.52|      0.65|      9.7|      5.0|[7.6,0.39,0.31,2....|      0.0|[479.531456024878....|[0.9590
6291204975...|      0.0|
|      7.9|      0.43|      0.21|      1.6|      0.106|      10.0|      37.0| 0.9966|3.17|      0.91|      9.5|      5.0|[7.9,0.43,0.21,1....|      0.0|[484.885809101067....|[0.9697
7161820213...|      0.0|
|      8.5|      0.49|      0.11|      2.3|      0.084|      9.0|      67.0| 0.9968|3.17|      0.53|      9.4|      5.0|[8.5,0.49,0.11,2....|      0.0|[491.163209482913....|[0.9823
2641896582...|      0.0|
|      6.9|      0.4|      0.14|      2.4|      0.085|      21.0|      40.0| 0.9968|3.43|      0.63|      9.7|      6.0|[6.9,0.4,0.14,2.4....|      1.0|[3.29191168408642....|[0.0065
8382336961...|      1.0|
-----
only showing top 5 rows

None
Test Accuracy of the Wine Quality Prediction Model: 0.9843627834245584
Weighted F1 score of the Wine Quality Prediction Model: 0.9779339666913879
```